

Model policy vs sweep-best policy — repeat measurements

Current model after early-diversity change. Gap = 100 × (model time / best time − 1).

Workload	Worker budget	Model allocated policy (modal if varying)	Model time (ms)	Sweep-best policy (remeasured)	Best time (ms)	Extra time (ms)	Gap from best (%)
large arrays	6	Dynamic (5, 5) P=1; order=LR; BF=1	5.4888	Dynamic (5, 3) P=1; order=LR; BF=1	5.6824	-0.1936	-3.4%
large arrays	8	Dynamic (7, 7) P=1; order=LR; BF=1 * 89% of native samples	6.7797	Dynamic (5, 3) P=1; order=RL; BF=1	5.5547	+1.2250	+22.1%
Mixed: compute	6	Dynamic (5, 5) P=1; order=LRabc; BF=0 * 89% of native samples	0.6893	Dynamic (5, 5) P=1; order=LabcR; BF=1	0.7027	-0.0133	-1.9%
Mixed: loop skew	6	Dynamic (5, 1) P=1; order=LRabc; BF=0	0.2734	Dynamic (5, 1) P=1; order=aLRbc; BF=1	0.2630	+0.0104	+4.0%
Mixed: short ordinary	6	Dynamic (1, 1) P=1; order=LRabc; BF=0	0.2254	Dynamic (5, 1) P=1; order=aRLbc; BF=1	0.2401	-0.0147	-6.1%

Policy: (left, right) loop widths; P = parent-worker limit; order = task priority; BF = backfill (1 on, 0 off).

L/R = loops; a/b/c = ordinary tasks. Priority order is not a serial execution sequence. Width 1 runs inline.

Times: median of process medians. Native time includes policy overhead. * = varying policies; time includes all native choices.

Best is the prior exhaustive-sweep winner, remeasured here. No new full sweep; negative gaps can reflect timing variation or a better native policy.